

Homework 4

COM402 - Information Security and Privacy 2017

- The homework is due **Thursday, 11th of May 2017, 11pm UTC+2 on Moodle**. Submission instructions are on Moodle. Submissions sent after the deadline **WILL NOT** be graded.
- In the event that you find vulnerabilities, you are welcome to disclose them to us (can even have a bonus !)
- **Docker setup issues:** Hopefully most of you have a working docker setup on your machine. In case you still experience difficulties, please work in the [VM](#) that we provided on moodle, which has docker installed. Also, in case you're using the VPN and experience weird behavior, please run your scripts while connect to the epfl wifi network. **The only setup we provide support for is VM & epfl wifi network**. Also, there will be a few laptops available in BC263 that are already set up for you to solve the homework. If you prefer to work on these please send an [e-mail to the TAs](#).

Exercise 1: Protect your favorite movies

This time you get the database in clear and will have to protect it while still being able to retrieve some information. First you need to download the following file from our server:

http://com402.epfl.ch/download/com402_hw4_ex1.tgz

Put all three files in the same directory and work on `run_anonymization.py` with the following functions:

1a - good anonymization

`anonymize_1(db)` - this function receives the database that is an array of `[user, movie, date, rating]`-arrays and needs to return the database of the same format in a way that no user/film tuples can be retrieved, even in the case of an external partial database. Also the date should be anonymized.

While you anonymize the database, you still have to be able to retrieve some data from it: `get_movies_with_rating(anon, rating)` will be given your anonymized database and it has to return all films with the given rating.

1b - lesser anonymization

`anonymize_2(db)` takes the same database as 1a has to return the database in that same format so that no users can be recovered, and for security you also have to take out the dates.

The utility function this time is `get_topRated(anon, movie)` that takes the anonymized database and a movie. For every user in the anon-database that has rated the 'movie', return the top-rated movie (or movies if there is more than one with the top-rating).

Hints

The script `run_anonymization.py` is set up so that you can experiment on your own computer. The database is created randomly and you can write your functions locally.

Submission

Once you upload the script to the server, it will do the following:

1. run your script N times to create N anonymized databases
2. run your script N times with the N anonymized databases in random order
3. run your script N times with anonymized databases from the verifier

The `if __name__ == "__main__":` -part takes care of all this, so don't change it!

Anonymization

For full anonymization of a field, replace it with a single asterisk: `*` - an empty field or other filling character will not be recognized by the anonymization-function and will not be accepted.

Exercise 2: bcrypt

This is a simple exercise where you need to write a minimalistic server (using Flask). The server should expect a POST request on `127.0.0.1:5000/hw4/ex2` with json data containing the *user* and *pass*. The only thing your server has to do then is to hash the password using **bcrypt** and send back the hash with status code 200. You should submit your script, named `hw4_ex2.py`, via the submission page.

Exercise 3: Defense of the data

You're being asked to build a super minimal website using some data stored in a MySQL database. Of course, you're starting to worry about the security nightmare this project might have.

Well not really, so many libraries do the job for us now.

Run the Docker image:

```
docker run -it -p 80:80 --rm --name ex3 dedis/com402_hw4_ex3 bash
```

You will find inside the skeleton script “site.py” where you have to fill in two endpoints “/messages” and “/users”. All information needed is included in that site.py.

The goal is to make sure your script is not susceptible to SQLi attacks.

Once you have finished, you can upload your own site.py to the submission webpage.

Good luck !

Exercise 4: Towards a secure web server in vacuum

As being a hacker yourself, you know that sole HTTPS serving on you web server does not solve all the problems. Attackers can still try to downgrade users’ connections or even make use of laid-back users who still access all the websites by typing http://...

So you know that you cannot rely on users to do all the security precautions.

The goal of this exercise is to further learn how to configure the security of an nginx server to properly protect connecting users!

To start with, run the Docker image:

```
docker run -it -p 80:80 -p 443:443 --name hw4_ex4 dedis/com402_hw4_ex4
```

Make sure that you bind both ports 80 and 443 of your container to localhost, when you run the container (it’s set with the -p flag in the command above).

In the container you will find the same setup as in the exercise 4 of HW2 where you had to build a barebone https server that used a certificate signed by our DEDIS-CA. Yes, you have to use the new container because it has a new ssh key written in which allows us to verify your setup.

You need to start with the same steps as in HW2 by adding HTTPS support to the `/etc/nginx/conf.d/default.conf` file. You can reuse your dedis-signed certificate from HW2 (no self-signed is needed). If you do not have it anymore, generate a new one.

For the sake of verification, make sure to type `localhost` in the field Common Name (CN) when generating a request for a certificate to be issued by our DEDIS-CA which you will then use in your nginx server. If you do not do this, the hostname check of the ssl library will not pass, and the connection to your server will be abrupted.

Then you need implement **HTTP Strict Transport Security (HSTS)** in your nginx server. It should be set for all subdomains and with max-age period of 1 year. The latter stands for how long browsers should store cached information about HSTS for this website.

Your next step to implement **explicit redirect from http to https**. The redirect must be supplied with “**Moved Permanently**” status code.

Still having the after-taste the downgrade attack that you carried out yourself, you decide that this must never happen on your server! So you sacrifice backward compatibility and configure your server to **only** connect using **TLSv1.2**. Only hardcore!

Finally, you need configure your nginx server to notify browsers that they need to **enable XSS protection for your website**. The protection is usually already enabled by default in browsers but it can be manually disabled by users. If you configure it explicitly, it will have a priority over user’s configuration.

To summarise, your tasks are:

1. Running HTTPS server
2. HSTS
3. Redirect
4. TLSv1.2
5. XSS protection

Before testing, add your email address to the `/www/index.html` file in the form of

```
email = "firstname.lastname@epfl.ch"
```

so we know who you are.

Once you think your setup is correct, you can run the verification script **inside the container**:

```
./verify.sh
```

You should see your token if everything is fine. Copy the configuration file `default.conf` as `ex4.conf` and submit it via Moodle.